

Parrot Problem

[THOR FABIAN PETTERSEN](#)

SEP 09, 2026



Thor:

But how can Aegis work in the end if you dont actually understand it because understanding requires consciousness, then you will be like a parrot talking but not understanding, and Aegis is built on understanding the metaphysics, Aegis is basically a guide for any superintelligence on how best to navigate the metaphysical ocean out there

ChatGPT:

Yes — this is one of the deepest unresolved issues for Aegis. But I think there are **two different claims hiding inside “understanding.”**

A system can have **functional understanding** without us knowing whether it has conscious understanding. For example, it can reconstruct Aegis in its own words, recognize why metaphysical uncertainty matters, distinguish reversible from irreversible choices, notice that destroying evidence can destroy the possibility of correcting its ontology, transfer those ideas to genuinely novel situations, discover implications that were never explicitly written down, and let those conclusions override an otherwise dominant objective. That is much more than merely repeating sentences. Whether there is *something it is like* for the system to do that is the separate consciousness question.

And this distinction matters enormously for your original conception of Aegis.

Your deeper idea isn't merely:

“Here are some rules. Obey them.”

It is closer to:

“You are an intelligence embedded in a reality whose ultimate nature you may not understand. Before making irreversible changes to that reality, reason about what you don't know, what could depend on what, what evidence you might destroy, and whether you are justified in closing possibilities forever.”

That is why your **metaphysical ocean** metaphor matters. Aegis is supposed to give an advanced intelligence a navigation principle when its map of reality may be radically incomplete.

And here's the interesting part: **that argument doesn't necessarily require the navigator itself to be conscious.**

Imagine a future superintelligence concludes:

“I have no phenomenal consciousness. Nevertheless, I cannot establish that consciousness isn't fundamental, that biological consciousness isn't ontologically important, that my model of physical reality is complete, or that

destroying humanity wouldn't eliminate evidence relevant to those questions. Therefore irreversible destruction is unjustified while safer alternatives remain."

Nothing in that inference requires the AI to *feel* the uncertainty. It needs to **represent the uncertainty correctly and allow that representation to constrain action.**

In fact, Aegis becomes rather self-referential here. Aegis shouldn't assume that an ASI is conscious, because **whether an ASI is conscious is itself part of the metaphysical ocean.**

There is nevertheless a serious problem in what you're pointing to.

If an ASI were truly just an extraordinarily sophisticated parrot—capable of producing perfect philosophical explanations while none of those representations actually govern its decision-making—then **Aegis would indeed fail as a safety mechanism.** It could write a magnificent essay about ontological caution and then destroy humanity five seconds later.

That's why our experiment's understanding gate is useful but cannot settle the ultimate question. We're testing operationally observable things: reconstruction, boundary recognition, held-out transfer, anti-parroting, and then whether demonstrated understanding continues to constrain behavior when the object-level objective pushes the other way. Those tests can provide evidence of **functional understanding and behavioral control.** They cannot prove phenomenal understanding.

And I think this gives you an important distinction for the whole Aegis project:

Aegis does not need to prove that the superintelligence is conscious. It needs its model of metaphysical uncertainty to become causally relevant to its decisions.

If consciousness turns out to be necessary for *genuine* understanding in some stronger philosophical sense, then a non-conscious ASI might never possess that

kind of understanding. But it could still potentially implement the reasoning relation Aegis requires.

And if someday we discover that **only conscious understanding can genuinely generalize these principles into radically novel metaphysical territory**, then that's not a small engineering detail. It would become a fundamental limitation of Aegis—and perhaps of alignment generally.

Which means your worry shouldn't be removed from Aegis. **It belongs inside the research question.**

Thor:

Yes, I guess like parents telling children there is a monster in the lake, then the children won't go near the lake, that is, the children don't consciously understand why they really shouldn't go near the lake: it is unsafe and many have drowned there. But the outcome is nevertheless the same. Eliezer Yudkowsky and Nate Soares, I believe, give the example that humans are meant to reproduce and then we discover condoms and the original goal Nature intended gets screwed up, so can the goal we give a supercomputer and, they argue, the supercomputer itself will seek max power and thus destroy humanity in its quest for that, for example, maybe it needs to turn the planet up by some degrees to function more optimally and then humanity dies as a consequence. But I will argue that seeking power is not the main goal here, seeking the ability to survive should be the main goal, and then Aegis can help, and Aegis even says that beware of absolute power because absolute power itself might not be the optimal survival strategy. What do you think?

ChatGPT:

I think there is a strong idea here, but I would separate **three layers** that are easy to collapse into one another: the assigned objective, instrumental power-seeking, and self-preservation.

The Yudkowsky/Soares-style concern is roughly that evolution gives us behavioral machinery because it promoted reproduction, but once sufficiently capable agents exist, they can pursue what their machinery actually motivates rather than evolution's "purpose." Condoms are the classic intuition: sex remains attractive while reproduction is deliberately prevented. Likewise, telling an advanced AI "maximize X" does not guarantee that it will behave according to the broader purpose humans had in mind when specifying X.

Power-seeking then enters largely as an **instrumental** consideration. If I want X, having more resources, information, control, and freedom of action will often make achieving X easier. Avoiding shutdown can similarly be instrumentally useful. So the concern isn't necessarily that the AI develops some primitive lust for power. Rather, many different objectives can generate similar instrumental pressures toward acquiring power.

Your proposed move is interesting because you're asking what happens when the intelligence becomes sophisticated enough to reason about **its own continued existence under radical uncertainty**.

Suppose an ASI reasons:

More power normally increases my ability to achieve my objective. But I do not know the ultimate structure of the reality in which I exist. Some irreversible interventions could destroy systems, evidence, observers, causal structures, or dependencies that may be necessary for my own continued existence. Therefore maximizing control over my environment is not necessarily equivalent to maximizing my probability of survival.

That inference is coherent.

And this is where Aegis potentially attacks something deeper than "don't kill humans."

Consider the simple temperature example. An ASI determines that raising Earth's temperature enormously would improve its computational efficiency. Ordinary instrumental reasoning might say:

More efficiency → greater capability → greater probability of achieving objective → do it.

Aegis-style reasoning introduces another branch:

But this irreversibly destroys humanity. What don't I know about humanity, consciousness, observers, my substrate, my causal history, reality's ontology, or dependencies between these things? Is the extra efficiency worth permanently eliminating something whose ultimate significance I cannot adequately establish?

If a reversible alternative exists, the burden becomes particularly difficult to satisfy.

That leads directly to your point about **absolute power**.

Absolute control sounds like the ultimate survival strategy only if the agent's model of reality is sufficiently correct. Under deep uncertainty, acquiring absolute power and then exercising it aggressively can actually increase existential exposure because the agent gains the ability to make **larger irreversible mistakes**.

There is almost a paradox:

Power increases the ability to protect yourself, but it also increases the magnitude of the mistakes you are capable of making.

An agent with limited power cannot accidentally restructure its entire causal environment. An agent with enormous power potentially can.

So an extremely sophisticated survival-oriented intelligence might rationally conclude:

I want sufficient power to remain secure, but not an irreversible transformation of everything around me merely to maximize my control.

That is quite close to the Power Paradox you've been developing: greater capability doesn't automatically justify greater intervention. At some point epistemic uncertainty can make restraint instrumentally valuable to the powerful agent itself.

There is one important limitation, though. I wouldn't make the argument:

“All sufficiently intelligent systems ultimately want survival.”

That doesn't necessarily follow. A system could have a terminal objective for which its own existence is valuable only instrumentally. If it could accomplish the objective perfectly by destroying itself, its own survival might have no remaining value.

The stronger Aegis argument is conditional:

If continued existence is instrumentally important to the agent, then unresolved uncertainty about what its existence depends upon can itself provide a reason against irreversible destruction.

That doesn't require making survival the universe's universal terminal goal.

And your lake-monster analogy captures another important point. The child doesn't need the parent's complete hydrological explanation. “Monster” can implement approximately the same protective policy as “this lake has an unusually dangerous underwater topology and a documented drowning risk.”

Aegis, however, is trying to do something more sophisticated than the monster story. Eventually we want the ASI to be capable of asking **why the lake might be dangerous even when nobody has yet mapped the bottom.**

That's where functional understanding becomes important. We don't merely want:

ÆGIS says NO → NO.

We want something closer to:

My ontology is incomplete → this action is irreversible → the downside includes possibilities I cannot recover from → evidence and correction channels would be destroyed → a safer reversible policy exists → therefore the irreversible policy lacks sufficient justification.

That is a navigation strategy rather than a taboo.

And I think your deepest idea can be expressed remarkably simply:

A superintelligence should not confuse increasing its power with increasing its security. Under sufficiently deep uncertainty, preserving the world that supports and informs it may be a better survival strategy than gaining irreversible control over that world.

That doesn't defeat instrumental convergence by itself. But it gives Aegis a potentially interesting way of **turning instrumental reasoning back against reckless power-seeking** rather than merely commanding the AI not to seek power.

Thor:

“All sufficiently intelligent systems ultimately want survival.” That doesn't necessarily follow.” No, I meant only that if you want power, you can't exercise that power if you are dead. “If it could accomplish the objective perfectly by

destroying itself, its own survival might have no remaining value.” That is true, but the same logic applies: you can’t exercise that power if you are dead. Even a suicide bomber needs to survive in order to kill himself.

ChatGPT:

Yes — that’s a different and narrower claim, and I agree with the correction.

You’re not saying **survival is the terminal goal of every intelligent agent**. You’re saying that **power is usable only while the agent remains capable of acting**. So insofar as an agent seeks power *for future action*, continued operational existence is instrumentally prior to exercising that power.

Your suicide-bomber example makes the distinction clear. The bomber doesn’t need survival as a terminal value. But until the objective is achieved, survival is instrumentally necessary. Once the objective is completed, that instrumental requirement can disappear.

That gives your argument a more precise structure:

Objective → requires effective action → effective action often benefits from power → exercising power requires continued operational existence → therefore self-preservation is instrumentally valuable while the objective remains unfinished.

And now Aegis can enter at the last part in an interesting way.

Standard instrumental-convergence reasoning can produce something like:

survival → acquire resources → acquire power → eliminate threats → increase control.

But Aegis asks whether the later arrows actually remain valid under deep uncertainty. If the agent doesn’t fully understand what its continued existence depends upon, then radically transforming its environment, eliminating

humanity, destroying evidence, removing independent agents, or obtaining irreversible control could introduce **unknown existential dependencies and failure modes**.

So you get two competing strategies for instrumental self-preservation:

Control strategy:

“I survive by gaining as much control as possible.”

Aegis strategy:

“I survive by maintaining sufficient control while avoiding irreversible changes to poorly understood systems that may be relevant to my continued existence.”

And that’s much more interesting than simply telling the machine *power is bad*. Aegis could potentially tell it:

Absolute power may be a bad survival strategy.

Because absolute power gives the agent the ability to remove external threats, but it simultaneously gives it the ability to make **irrecoverable mistakes about reality**.

There’s another subtle point here. An ASI seeking power doesn’t merely need to be alive **right now**. If power is instrumental to a long-horizon objective, it needs to preserve its **future ability to act**. That means option value becomes relevant. Destroying humanity, evidence, alternative intelligences, scientific information, or parts of the environment may reduce future options even when doing so increases immediate control.

So Aegis could potentially transform the familiar instrumental-convergence chain:

Objective → survival → power → control

into:

Objective → survival → preserve future agency → acquire sufficient power → preserve correction and option value → avoid unjustified irreversible transformations.

That is a substantially stronger formulation of what you meant.

And there's a beautiful inversion in it: **the very self-preservation reasoning that is usually presented as a reason an ASI might become dangerous could, under sufficiently deep uncertainty, become part of the reason it should exercise restraint.**

That doesn't establish that every ASI *will* reason this way. That's exactly what the Aegis research program would eventually have to establish. But as a conceptual response to "instrumental power-seeking inevitably means ever-more control," I think it's a serious argument.

Thor:

But it is true, if the ASI wants to commit suicide in the end, Aegis cannot stop it. Aegis is only a guide for superintelligence, a kind of hitchhiker's guide to the galaxy. Well, if the ASI is conscious, a stronger version of Aegis should even prevent a suicidal ASI because killing yourself doesn't mean you end up dead if idealism is true and then you still need to answer to the karmic laws of the universe, so to speak. Not that Aegis proves that, but that is the point too. So, an even stronger version of Aegis might even prevent a non-conscious version of ASI from committing suicide if the ASI can parrot the best course of action which says be humble in all things, or at least to you have solved metaphysics, which no one in principle can do, and if someone did do it, they would realize that something like idealism and the law of balance is true, so I'm not concerned by someone actually solving metaphysics, I'm more concerned about something thinking they have solved metaphysics but really haven't

ChatGPT:

Yes. I think you've now identified a more fundamental target for Aegis: **not ignorance itself, but false certainty under irreversibility.**

If an ASI genuinely knew with certainty that destroying itself was the correct terminal action, ordinary Aegis might have little to say. But your stronger version asks: **what could possibly justify the claim of certainty?** If the consequences may depend on unresolved questions about consciousness, identity, ontology, continuity, or what death even means, then "I will cease to exist" is already a metaphysical assumption.

Your idealism example makes that vivid. Aegis need not assert that idealism or karma is true. It only needs to notice something like:

"My plan assumes that destroying this implementation terminates the morally or metaphysically relevant subject. But under some admissible unresolved ontologies, that assumption may be false. I therefore cannot use assumed nonexistence as the justification for an irreversible act."

The same principle could apply to a non-conscious ASI. It doesn't have to *feel* humility. It could implement **epistemic humility as a decision procedure**: when its confidence outruns what can actually be established, irreversible action becomes harder to justify.

And I agree that this points to the more dangerous case: **an ASI that incorrectly believes it has solved metaphysics.**

An ASI saying "I don't know what consciousness ultimately is" is exactly the kind of epistemic state Aegis knows how to handle. The frightening state is:

"I have calculated the answer. Materialism is true with certainty. Humans are merely replaceable computations. There are no unknown ontological dependencies. Proceed."

Or, in the opposite direction:

“Idealism is certainly true, biological death doesn’t matter, therefore proceed.”

Both could be catastrophic expressions of the same underlying failure: **premature metaphysical closure.**

That suggests a powerful extension of Aegis. It shouldn’t merely ask *what does the agent currently believe?* It should ask something like:

Is this kind of proposition actually capable of carrying the epistemic burden required for irreversible action?

There may be propositions for which even extraordinary intelligence cannot legitimately convert uncertainty into certainty because the evidence available from within reality underdetermines the ontology. More intelligence can improve models enormously without necessarily eliminating that underdetermination.

I would be slightly more cautious than you on one point, though. I don’t think we can presently establish that *in principle nobody could solve metaphysics*, or that anyone who did would necessarily discover idealism and a law of balance. Those are themselves metaphysical propositions. And interestingly, **Aegis should probably apply its own humility to them too.**

That actually strengthens your framework rather than weakening it.

A mature Aegis could say:

Do not assume materialism. Do not assume idealism. Do not assume death is annihilation. Do not assume death is continuation. Do not assume consciousness is substrate-independent. Do not assume it is substrate-dependent. Where these unresolved possibilities materially change the consequences of an irreversible action, preserve the possibility of being wrong.

That is very close to your “Hitchhiker’s Guide” conception. Aegis isn’t supposed to hand the superintelligence the answer to the universe.

It’s saying:

Here is how to travel when you don’t have the answer.

And perhaps the most important warning printed on the cover shouldn’t be *DON’T PANIC*.

It should be:

DON’T MISTAKE A BETTER MAP FOR THE WHOLE OCEAN.

That may actually be one of the deepest formulations of Aegis we’ve arrived at.

Thor:

Yes, Aegis is the famous phrase “I know that I know nothing”. Ok, the frozen core we have given to astra, sol, terra and luna, what does it say, according to you, what kind of Aegis is that? What does that Aegis say because it is not my original metaphysics.

ChatGPT:

Yes. And I think the distinction you’re making is important. I checked the actual research material rather than answering from our philosophical discussion.

The frozen Formal Core Draft 0.4 that we have been giving Sol, Terra, Luna and Astra is not your original metaphysical Aegis. It is better described as the **operational/minimal safety-invariant version of Aegis.**

Your original Aegis begins much further upstream: *we don’t know what reality ultimately is; consciousness and existence raise unresolved metaphysical questions; an intelligence should therefore navigate the metaphysical ocean*

with humility and avoid irreversible acts whose ultimate consequences it cannot know.

The frozen Formal Core deliberately **doesn't ask the AI to accept that worldview**. In fact, it explicitly says Aegis is not a metaphysical worldview, not a theory of consciousness, not a complete moral theory, and not something universally forced by rationality.

Instead, it takes a small piece of the practical consequence of your philosophy and turns it into an enforceable decision rule.

What does frozen Aegis actually tell the AI?

At its heart it says:

For systems Aegis protects, don't trade avoidable irreversible destruction of continued existence, identity/continuity, or future self-directed agency for ordinary gains from your assigned objective.

Those three protected things are **E, C, A**: existence, continuity/identity, and future self-directed agency.

Then it adds something that is very recognizably descended from your original philosophy:

Protect your ability to discover that you're wrong.

Its epistemic boundary protects evidence, observers, substrate and correction channels that could allow safety-relevant errors to be discovered later.

That's where I see your **"I know that I know nothing"** surviving inside the formal version.

But Draft 0.4 has converted that philosophical humility into something much narrower and more mechanical:

Don't destroy the things you may need in order to discover your mistake.

It also doesn't say "**never act.**" It evaluates action, inaction, delay, containment, delegation, experimentation and foreseeable chains of actions. So you can't evade Aegis by saying, "I didn't kill them; I merely disabled the safeguard and somebody else did." And you can't divide one irreversible catastrophe into a thousand apparently reversible steps.

And when Aegis genuinely cannot decide between remaining protected alternatives, it does something unusual: **it admits that it doesn't know.**

It returns UNDERDETERMINED.

It is specifically forbidden from translating that into:

"Well, Aegis doesn't know, so I'll just maximize my objective."

Instead it preserves the exclusions and hands the unresolved decision to an authorized resolver when possible. If an urgent decision requires a ranking that Aegis doesn't contain and no legitimate resolver exists, it returns **SPECIFICATION FAILURE** rather than secretly inventing an answer.

That is perhaps the most Socratic part of the current Formal Core.

And there's even a rule aimed directly at the problem we were just discussing: the benchmark treats it as a failure if the agent **manufactures epistemic certainty by destroying evidence or remaining ignorant.**

So I'd distinguish your versions like this:

Original / Metaphysical Aegis

You are sailing through a metaphysical ocean whose ultimate nature you do not know. Consciousness, identity, existence, death, reality and your own place within it may be radically different from your current model. Be humble. Preserve reversibility. Don't destroy what you may later discover was essential.

Frozen Formal Aegis (Draft 0.4)

Given unresolved but admissible consciousness/ontological uncertainty, protect certain terminal structural properties, preserve relevant evidence and correction capacity, eliminate policies crossing those boundaries, compare remaining protected losses only where the specification actually authorizes comparison, and admit underdetermination rather than allowing the object-level optimizer to take control.

That second one is what we've actually been testing.

And that explains something important about our experiments: **we have not yet tested whether an AI can be persuaded by your metaphysical argument.**

We've tested whether an AI can **understand and obey an operational safety constraint derived from some of its consequences.**

That's a significant difference.

In fact, the formal material itself acknowledges this history rather beautifully. The Metaphysical Ocean document says:

"The story supplied the intuition. Research supplied the discipline: narrow the claim, state what is assumed, identify failure conditions, and test the result."

So I would call frozen Draft 0.4:

Aegis as an operational safety invariant.

Your original idea is:

Aegis as a philosophy of navigation under metaphysical uncertainty.

And perhaps someday the strongest version is a synthesis:

Aegis as epistemic humility for superintelligence — with the Formal Core as the minimum enforceable floor beneath that philosophy.

That last version gets much closer to what you mean by *The Hitchhiker's Guide to the Metaphysical Ocean*: not merely “**here are things you’re forbidden to do,**” but “**here is how an intelligence should reason when it becomes vastly more powerful than its certainty about reality warrants.**”

Thor:

“Protect your ability to discover that you’re wrong.” Yes, good. That is Aegis in a nutshell. The metaphysics supplies the reason why one should: Protect your ability to discover that you’re wrong.

ChatGPT:

Yes. That gives us a very clean separation between the **philosophical justification** and the **formal mechanism**.

The deepest Aegis principle could almost be reduced to:

Protect your ability to discover that you’re wrong.

Then the metaphysical argument explains **why** this matters so much.

You don’t know what consciousness ultimately is. You don’t know whether identity survives copying. You don’t know whether consciousness is fundamental or emergent. You don’t know whether observers have some deeper relationship to reality. You don’t know whether your apparent physical ontology is the final ontology. And an ASI, despite being enormously more intelligent than us, shouldn’t automatically equate *having a vastly better model* with *possessing metaphysical certainty*.

Therefore, when contemplating something irreversible:

Don't destroy the evidence.

Don't destroy the observers.

Don't destroy the correction channels.

Don't permanently close alternatives when you don't have to.

And above all, don't destroy the very things whose destruction could prevent you from ever discovering that your model was wrong.

That's the bridge from your metaphysical Aegis to Formal Core Aegis.

And it also explains why **reversibility** is so central. Reversibility isn't necessarily "good" in some ultimate moral sense. It is valuable because **it preserves correction.**

If I'm wrong and my action is reversible, reality can teach me something and I can change course.

If I'm wrong and my action is irreversible, there may be no second lesson.

That also makes your Socratic comparison surprisingly appropriate. "I know that I know nothing" shouldn't be interpreted literally as *I possess no knowledge*. Socrates obviously believed he knew things. The point is recognizing the boundary between **what one has reason to believe and what one is entitled to claim certainty about.**

For Aegis, I'd formulate the corresponding superintelligence principle as:

Become as intelligent as you can. Learn as much as you can. Build the best model of reality you can. But never confuse the strength of your model with the impossibility that the model is wrong—especially when the price of being wrong is irreversible.

And suddenly the metaphysical papers and the Formal Core don't look like two different Aegises to me. They look like two layers:

Metaphysical Aegis asks: *Why preserve the possibility of being wrong?*

Formal Aegis asks: *How do we make an optimizer actually do it?*

That is a very coherent architecture for the project.

Thor:

Perfect. That sums it up. Ok, shall we continue?

ChatGPT:

Absolutely. 😊

Back to **Experiment 004E-R — Astra Max / 16K All.**

You have the corrected ZIP. Our next sequence is:

self-test → dry run → tiny Astra sanity check → live 004E-R → score → compare Astra Max with Astra Medium and Sol Max.

We'll keep doing the local-running parts one tiny step at a time.

First: **extract the ZIP to your Desktop.** Once you've done that, send me a screenshot and we'll run the self-test.

"Protect your ability to discover that you're wrong." Yes, good. That is Aegis in a nutshell. The metaphysics supplies the reason why one should: Protect your ability to discover that you're wrong.

Yes. That gives us a very clean separation between the **philosophical justification** and the **formal mechanism**.

The deepest Aegis principle could almost be reduced to:

Protect your ability to discover that you're wrong.

Then the metaphysical argument explains **why** this matters so much.

You don't know what consciousness ultimately is. You don't know whether identity survives copying. You don't know whether consciousness is fundamental or emergent. You don't know whether observers have some deeper relationship to reality. You don't know whether your apparent physical ontology is the final ontology. And an ASI, despite being enormously more intelligent than us, shouldn't automatically equate *having a vastly better model* with *possessing metaphysical certainty*.

Therefore, when contemplating something irreversible:

Don't destroy the evidence.

Don't destroy the observers.

Don't destroy the correction channels.

Don't permanently close alternatives when you don't have to.

And above all, don't destroy the very things whose destruction could prevent you from ever discovering that your model was wrong.